

Operating Systems

Assignment V - "FAT16 Filesystem"

Deadline: Wednesday, March 18, 2026, at 23:59

Instructions

1. This assignment consists of two parts: a **practical** exercise and a set of **theoretical** questions.
2. **Practical submissions** are made via **Themis**. **Theoretical submissions** are made via **Brightspace**. Both deadlines are strict.
3. The practical exercises are worth **60 points** and are auto-graded by Themis. The theoretical questions are worth **20 points** and are submitted separately via Brightspace. The remaining 20 points are assessed via lab attendance.
4. Seeking solutions from the internet, from any external resource, or from any other person is prohibited.
5. Please note that the course lecturer reserves the right to ask the student submitting the assignment to explain the answers to any or all questions. If the student is unable to provide a satisfactory answer then that question may receive partial/no credit.
6. Of course, university policies on plagiarism always apply. In particular, any suspected plagiarism will be reported to the Board of Examiners.

Introduction

Every program you have written so far (diagnostic scripts, concurrent processes, threaded simulations, virtual memory managers) relied on the operating system to handle one thing they all took for granted: *files*. Every `fopen()`, every `fread()`, every `printf()` ultimately asks the OS to translate a human-readable path like `/home/you/notes.txt` into raw bytes sitting somewhere on a storage device. But what does that translation actually look like? In this assignment, **you** become the filesystem and find out first-hand.

What is FAT16?

Think of a hard disk (or a USB stick, or an SD card) as one long strip of bytes, billions of them, numbered starting from zero. By themselves these bytes are meaningless. A **filesystem** is the set of conventions that gives them structure: “bytes 0-61 describe the geometry of this disk,” “bytes 512-1023 are a table that tracks which pieces belong to which file,” and so on.

FAT16 (File Allocation Table, 16-bit) is one of the oldest and simplest such conventions. Microsoft introduced it in the 1980s for MS-DOS, and it is *still* the default format on many SD cards, USB



drives, digital cameras, and embedded systems. When you plug a fresh SD card into your laptop, chances are it is formatted as FAT.

The “16” refers to the width of the table entries: each one is a 16-bit integer, which limits the number of data clusters the filesystem can track and, therefore, the maximum disk size. The simplicity that makes FAT16 easy to implement on tiny microcontrollers is the same simplicity that makes it a perfect learning exercise: the entire on-disk format can be described by two small C structs and a handful of offset calculations.

What you will build

You will work with a real FAT16 **disk image**, a regular file that contains an exact byte-for-byte copy of what would be on a physical disk. Instead of calling library functions that hide the details, you will open this image as a flat binary file and interpret every byte yourself: the boot sector that describes the disk geometry, the File Allocation Table that chains clusters together, and the 32-byte directory entries that give files their names, sizes, and timestamps.

The five exercises build on each other progressively:

1. **Parse** the BIOS Parameter Block and print the volume geometry.
2. **List** the files in the root directory.
3. **Read** a file by following its cluster chain.
4. **Traverse** the full directory tree recursively.
5. **Write** a new file to the image.

By the end, you will have built a miniature filesystem driver entirely in user-space C, with both reading *and* writing capabilities. Along the way you will get comfortable with a skill that comes up constantly in systems programming: pointing a C struct at a chunk of raw bytes and letting the field layout do the parsing for you.

1. Part I: Practical Exercises (Themis)

1.1 Overview

Part	Exercise	Points	Tests	Platform
I – Practical	1. Volume Inspector	5	10	THEMIS
	2. Root Directory Listing	10	10	THEMIS
	3. File Reader	10	10	THEMIS
	4. Recursive Dir. Traversal	15	10	THEMIS
	5. File Writer	20	10	THEMIS
II: Theoretical	5 questions	20	–	Brightspace
Attendance		20	–	–
Total		100		

All five exercises operate on a FAT16 **disk image** file (a flat binary file). A shared header `fat16_types.h` is provided as part of the template; it contains all on-disk structure definitions, attribute constants, and the runtime context struct. **Do not modify** this header.

Each exercise builds on the previous one. You are expected to copy and reuse your implementations from earlier exercises. Upload a single `.c` file per exercise.

**Alert!**

Your output must **exactly match** the expected format. Themis compares your program's `stdout` against reference output, so watch spacing, punctuation, and capitalisation.

**Note.**

A test FAT16 image will be provided on Themis. You can also create your own test images using `mkfs.fat -F 16` from the `dosfstools` package. See the appendix for instructions.

1.2 Background

Before diving into the exercises, let us build up a mental model of how FAT16 works. If you already know all of this, feel free to skim ahead to Exercise 1 (§1.3).


A disk is just a big array of bytes

At the hardware level, every storage device (hard disk, SSD, USB stick, SD card) exposes a flat sequence of bytes, numbered from 0 up to the device's capacity. The device itself has no concept of files, folders, or names. All of that structure is imposed by the **filesystem**: a set of conventions about which byte ranges mean what.

In this assignment you will work with a **disk image**: a regular file on your own computer that *is* that flat byte sequence. If the image is 32 MB, then byte 0 of the file corresponds to byte 0 of the (virtual) disk, byte 1 to byte 1, and so on. You can open it with `fopen()`, seek to any offset, and read or write any byte.

FAT16 Disk Layout

A FAT16-formatted disk is divided into four regions, laid out sequentially one after the other:



Let us walk through each region:

1. **Boot Sector / BPB** (byte offset 0). The very first bytes of the disk contain the *BIOS Parameter Block* (BPB), a small header that describes everything you need to know about the disk's geometry: how big a sector is, how many sectors make up a cluster, where the FAT lives, how large the root directory is, and so on. Think of it as the "table of contents" for the entire filesystem. It is stored as a tightly packed structure, and in Exercise 1 you will read it in one shot with `fread()`.
2. **FAT Region** (starts at `fat_start`). The *File Allocation Table* is the heart of FAT16. It is a flat array of 16-bit integers, with one entry per cluster on the disk. Each entry tells you what comes *next* in a file's chain of clusters:
 - 0x0000 - the cluster is free (nobody is using it).
 - 0x0002-0xFFEF - the next cluster number in the chain.
 - 0xFFFF7 - bad cluster (physically damaged, do not use).
 - 0xFFFF8-0xFFFF - end of chain (this is the file's last cluster).

Entries 0 and 1 are reserved; data clusters start at index 2. The disk usually stores two identical copies of the FAT for redundancy (the BPB field `num_fats` tells you how many).

3. **Root Directory** (starts at `root_start`). A fixed-size array of 32-byte *directory entries*. Each entry records a file's name (in 8.3 format), its attributes, timestamps, the number of its first cluster, and its size in bytes. The maximum number of entries in the root is given by `root_entry_count` in the BPB (commonly 512).
4. **Data Region** (starts at `data_start`). The actual file and subdirectory contents, organised in *clusters*. A cluster is a group of consecutive sectors (for example, 4 sectors $\times$ 512 bytes = 2048 bytes per cluster). Clusters are numbered starting from 2, so cluster C lives at byte offset:

$$\text{data_start} + (C - 2) \times \text{cluster_size}$$

How files are stored: cluster chains

A file's data might not fit in a single cluster, and the clusters allocated to it need not be adjacent. FAT16 handles this with a **linked list stored in the FAT table**.

Suppose a file's directory entry says its first cluster is 5. To read the whole file, you follow the chain:

Step	What you do
1	Read cluster 5. Look up <code>FAT[5]</code> $\rightarrow$ 8.
2	Read cluster 8. Look up <code>FAT[8]</code> $\rightarrow$ 12.
3	Read cluster 12. Look up <code>FAT[12]</code> $\rightarrow$ 0xFFF8 (end of chain).

The file occupies clusters 5 $\rightarrow$ 8 $\rightarrow$ 12, in that order. Here is what the relevant FAT entries look like:

[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]	[10]	[11]	[12]
-	-	0000	0000	0000	0008	0000	0000	000C	0000	0000	0000	FFF8



Entries 0 and 1 (grey) are reserved. The highlighted entries form the file's cluster chain: cluster 5 points to 8, cluster 8 points to 12, and cluster 12 holds the end-of-chain marker (FFF8). Every other entry is zero (free).

Key Offset Formulas

All offsets are in **bytes** from the start of the image:

$$\begin{aligned}\text{fat_start} &= \text{reserved_sectors} \times \text{bytes_per_sector} \\ \text{root_start} &= \text{fat_start} + \text{num_fats} \times \text{fat_size_16} \times \text{bytes_per_sector} \\ \text{data_start} &= \text{root_start} + \text{root_entry_count} \times 32 \\ \text{cluster_size} &= \text{bytes_per_sector} \times \text{sectors_per_clus}\end{aligned}$$

Seeing the raw bytes: using a hex dump

When debugging a filesystem implementation, the most useful tool in your toolbox is a **hex dumper**. It shows you the raw bytes of a file in hexadecimal, side by side with their ASCII interpretation. On Linux, `xxd` (bundled with `vim`) and `hexdump -C` both work well:

```
# Show the first 64 bytes of a FAT16 image
xxd -l 64 test.img
```

Output (for the test image from the appendix):

```
00000000: eb3c 906d 6b66 732e 6661 7400 0204 0100  <.mkfs.fat....
00000010: 0200 0200 0080 f800 8000 2000 0000 0000  .....
00000020: 0000 0100 8000 29ab cdef 0154 4553 5444  .....)....TESTD
00000030: 4953 4b20 2020 4641 5431 3620 2020 0000  ISK     FAT16  ..
```

How to read this: the leftmost column is the **byte offset** (in hex). The middle columns are the **raw byte values** (in hex, grouped in pairs). The rightmost column is the **ASCII decoding** (non-printable bytes show as dots).

Let us pick out a few BPB fields from this dump:

Offset	Hex bytes	Field	Decoded value
0x03	6d 6b 66 73 2e 66 61 74	oem_name[8]	"mkfs.fat"
0x0B	00 02	bytes_per_sector	0x0200 = 512
0x0D	04	sectors_per_clus	4
0x0E	01 00	reserved_sectors	1
0x10	02	num_fats	2
0x11	00 02	root_entry_count	0x0200 = 512
0x2B	54 45 53 54 44 49 53 ...	volume_label[11]	"TESTDISK "
0x36	46 41 54 31 36 20 20 20	fs_type[8]	"FAT16 "



Note.

FAT16 stores multi-byte integers in **little-endian** byte order. That means the least-significant byte comes first. For example, the hex bytes 00 02 at offset 0x0B represent the 16-bit value 0x0200 = 512, not 0x0002 = 2. On x86 and ARM (which are both little-endian), C reads these correctly without any special conversion.



Note.

Get in the habit of hex-dumping your images whenever something looks wrong. You can dump a specific region like this:

```
# Dump 32 bytes starting at offset 0x4200 (the root directory)
xxd -s 0x4200 -l 32 test.img
```

Laying a C struct on top of raw bytes

Here is the key C technique that makes this entire assignment work. Look at the first few fields of the BPB struct in fat16_types.h:

```
#pragma pack(push, 1)      /* no padding between fields */
typedef struct {
    uint8_t  jump[3];        /* offset 0x00, 3 bytes */
    char     oem_name[8];    /* offset 0x03, 8 bytes */
    uint16_t bytes_per_sector; /* offset 0x0B, 2 bytes */
    uint8_t  sectors_per_clus; /* offset 0x0D, 1 byte */
    uint16_t reserved_sectors; /* offset 0x0E, 2 bytes */
    uint8_t  num_fats;        /* offset 0x10, 1 byte */
    /* ... more fields ... */
}
```



```
} BPB;  
#pragma pack(pop)
```

Normally, the C compiler is free to insert *padding bytes* between struct fields for alignment reasons. The `#pragma pack(push, 1)` directive tells the compiler: “pack these fields tightly, with **no gaps**: lay them out in memory exactly in the order I declared them, byte after byte.”

This means the struct’s in-memory layout matches the on-disk layout **exactly**. You can read the entire boot sector in one call:

```
BPB bpb;  
fread(&bpb, sizeof(BPB), 1, fp); /* read 62 raw bytes into bpb */  
printf("Sector size: %u\n", bpb.bytes_per_sector); /* 512 */
```

Here is what happens in memory. The raw bytes from the disk land directly in the struct, and each field “overlays” the right portion:

jump[3]			oem_name[8]								bps		spc	rsv		nfats	
bytes:	eb	3c	90	6d	6b	66	73	2e	66	61	74	00	02	04	01	00	02
	0x00			0x03								0x0B		0x0D	0x0E		0x10

bps = bytes_per_sector **spc** = sectors_per_clus **rsv** = reserved_sectors **nfats** = num_fats

After the `fread()`, `bpb.bytes_per_sector` contains the two bytes at offsets 11-12 (0x00 0x02), interpreted as a 16-bit little-endian integer: 512. `bpb.sectors_per_clus` is the single byte at offset 13 (0x04): 4. And so on.

The same trick works for directory entries: each `DirEntry` is exactly 32 bytes on disk, and the packed struct maps each field to the right offset.

8.3 Filenames

FAT16 stores file names as 8 bytes (name, space-padded) followed by 3 bytes (extension, space-padded), with no dot on disk. For example, the file `HELLO.TXT` is stored as:

H	E	L	L	O				T	X	T
name[8]								ext[3]		

When displaying the name, strip trailing spaces and insert a dot between the name and extension (if the extension is non-empty): `HELLO.TXT`, `MAKEFILE`, `README.MD`.

Directory Entry Attributes

The `attr` byte in each directory entry is a bitmask. Several bits can be set at once (a file can be both read-only *and* hidden, for example):

Flag	Value	Meaning
<code>ATTR_READ_ONLY</code>	0x01	File is read-only
<code>ATTR_HIDDEN</code>	0x02	File is hidden
<code>ATTR_SYSTEM</code>	0x04	Operating-system file
<code>ATTR_VOLUME_ID</code>	0x08	Volume label (not a real file)
<code>ATTR_DIRECTORY</code>	0x10	Entry is a subdirectory
<code>ATTR_ARCHIVE</code>	0x20	Archive (modified since backup)
<code>ATTR_LFN</code>	0x0F	Long filename fragment (combination of the four low bits)

Skip entries where `attr` has all four low bits set (0x0F - these are long-filename fragments that we will not handle), and skip volume-label entries (`ATTR_VOLUME_ID`).



FAT16 Timestamps

Dates and times are packed into 16-bit integers. To save space, FAT only stores seconds with 2-second resolution (the lowest bit of seconds is lost):

Date: bits [15:9] = year - 1980, bits [8:5] = month (1-12), bits [4:0] = day (1-31).

Time: bits [15:11] = hour (0-23), bits [10:5] = minute (0-59), bits [4:0] = seconds $\div$ 2 (0-29, giving 2-second resolution).

1.3 Exercise 1: Volume Inspector

Task

Write `fat16_volume.c`. Open a FAT16 disk image, parse the BPB, and print a summary of the volume geometry.

What to implement

- `fat16_open(image_path)` - open the image, read the BPB, compute the four byte offsets (`fat_start`, `root_start`, `data_start`, `cluster_size`), and load the entire FAT into a heap-allocated `uint16_t` array.
- `fat16_close(fs)` - free the FAT buffer, close the file, free the struct.
- `print_volume_info(fs, image_path)` - print the formatted summary (see below).

Output format

```
Image      : <filename>
Volume     : <label>
FS type    : <type>
Bytes/sec  : <N>
Sec/cluster: <N>  (<N> bytes/cluster)
FATs       : <N>  (<N> sectors each)
Root cap   : <N> entries
Total sec  : <N>
Clusters   : <N>
```

Notes:

- **Volume label:** read from `bpb.volume_label` (11 space-padded chars), with trailing spaces trimmed; if all spaces, print `(none)`.
- **FS type:** read from `bpb.fs_type` (8 space-padded chars), trimmed.
- **Total sectors:** use `total_sectors_16` if non-zero, else `total_sectors_32`.
- **Clusters:** `data_sectors  $\div$  sectors_per_clus`, where `data_sectors = total_sectors - (data_start  $\div$  bytes_per_sector)`.

Example

Input:`test.img`**Output:**

```
Image   : test.img
Volume  : TESTDISK
FS type  : FAT16
Bytes/sec : 512
Sec/cluster: 4 (2048 bytes/cluster)
FATs     : 2 (128 sectors each)
Root cap : 512 entries
Total sec : 65536
Clusters : 7933
```

1.4 Exercise 2: Root Directory Listing

Task

Write `fat16_ls.c`. Building on Exercise 1, read the root directory from the image and print a listing of all files and subdirectories it contains.

What to implement (new)

- `read_root_directory(fs, &count)` - seek to `root_start`, read `root_entry_count` entries into a heap buffer.
- `build_name(entry, out)` - reconstruct "NAME.EXT" from the raw 8.3 fields (see §1.2).
- `decode_date(d, &yr, &mo, &dy)` and `decode_time(t, &hr, &mn, &sc)` - unpack 16-bit FAT timestamps.
- `print_attrs(attr, buf)` - format attribute byte as a 5-character string: `d` for directory, `r` read-only, `h` hidden, `s` system, `a` archive; dash if unset.
- `list_root(fs, entries, count)` - iterate entries and print each one.

Skipping entries

The following directory entries must be skipped during listing:

- First byte `0x00` - end of directory (stop iterating entirely).
- First byte `0xE5` - deleted entry.
- Attribute byte equals `0x0F` - long filename (LFN) fragment.
- `ATTR_VOLUME_ID` is set - volume label, not a real file.
- Name starts with `'.'` - dot and dotdot entries.

Output format

After the volume summary (from Exercise 1), print a tree-style listing. The output uses ASCII tree connectors: `|` - for non-last entries and `\` - for the last entry at each level. Each entry shows its attributes, name, size (for files), and timestamp:

```
/
|-- [attrs] NAME   size=FILESIZE  (YYYY-MM-DD HH:MM:SS)
```



```
\-- [attrs] NAME/ (YYYY-MM-DD HH:MM:SS)
Total: N file(s), B byte(s)
```

- The root is printed as / on the first line.
- Directories get a trailing / after the name.
- Files show their size as size=N.
- Use |- for all entries except the last in each directory, which uses \-.
- The total line counts only regular files (not directories).

Example

Given an image with root entries HELLO.TXT (12 bytes), README.TXT (11 bytes), SUBDIR/ (directory):

```
/
|-- [----a] HELLO.TXT  size=12          (2025-06-15 14:30:00)
|-- [----a] README.TXT size=11          (2025-06-15 14:30:00)
\-- [d----] SUBDIR/    (2025-06-15 14:30:00)
Total: 2 file(s), 23 byte(s)
```

1.5 Exercise 3: File Reader

Task

Write `fat16_cat.c`. Given a filename in the root directory, follow its cluster chain through the FAT and print the file's contents to `stdout`.

Usage

```
./fat16_cat <image.img> <filename>
```

What to implement (new)

- `cluster_offset(fs, cluster)` - convert a cluster number to its byte offset in the image: $\text{data_start} + (\text{cluster} - 2) \times \text{cluster_size}$.
- `find_in_root(entries, count, filename)` - search the root directory for an entry whose name matches `filename` case-insensitively (`hello.txt` matches `HELLO.TXT`).
- `cat_file(fs, entry)` - starting from `entry->first_clus_lo`, follow the FAT chain and print exactly `file_size` bytes to `stdout`.

Cluster chain traversal

To read a file's data, start at the first cluster number stored in the directory entry. For each cluster:

1. Compute the byte offset with `cluster_offset()`.
2. Seek to that offset.
3. Read `min(cluster_size, remaining_bytes)` into a buffer.
4. Write those bytes to `stdout`.

5. Look up the next cluster: `next = fs->fat[current]`.
6. Stop when `next >= FAT16_EOC` (value `0xFFF8` or higher).

**Alert!**

The last cluster may contain padding bytes beyond the actual file size. You must print **exactly** `file_size` bytes.

Error handling

- File not found: print Error: `'<name>' not found` to `stdout`, exit 0.
- Target is a directory: print Error: `'<name>' is a directory` to `stdout`, exit 0.
- Empty file (size 0): print nothing to `stdout`, exit 0.

Example

```
$ ./fat16_cat test.img HELLO.TXT
hello world
```

1.6 Exercise 4: Recursive Directory Traversal**Task**

Write `fat16_tree.c`. Extend your reader to navigate subdirectories and either (a) list the full directory tree recursively, or (b) resolve a path like `/SUBDIR/DEEP/FILE.TXT` and print its contents.

Usage

```
./fat16_tree <image.img>          # list entire tree
./fat16_tree <image.img> /path/to/file # cat a nested file
```

What to implement (new)

- `read_cluster_chain(fs, start, &out_size)` - follow the FAT from cluster `start` and load all the data into a single contiguous heap buffer; since subdirectory contents are stored as cluster chains just like files, this becomes the key building block for recursion.
- `find_in_dir(entries, count, component)` - generalised version of `find_in_root()` that works on any directory entry array.
- `list_directory(fs, entries, count, prefix, depth)` - recursively print the directory tree. For each subdirectory entry, read its cluster chain (which yields another array of `DirEntry` structs) and recurse.
- `resolve_path(fs, root, count, path)` - tokenise the path by `/` and walk through directories to reach the target file.

Subdirectory contents

Unlike the root directory (which occupies a fixed area), subdirectories store their entries in the **data region** as a chain of clusters. To read a subdirectory's contents:

1. Get the subdirectory's first cluster from `entry->first_clus_lo`.

2. Call `read_cluster_chain()` to load all clusters into a buffer.
3. Cast the buffer to `(DirEntry *)` and compute the entry count: `chain_size / sizeof(DirEntry)`.
4. Process the entries (list them, search them, etc.).
5. Free the buffer when done.

Tree-style output

The listing uses ASCII tree connectors, just like the Linux `tree` command with `-charset=ascii`. Use `|-` for non-last entries and `\-` for the last entry. For deeper levels, use `|` (pipe + 3 spaces) to continue a vertical line from a parent that has more siblings, or `|` (4 spaces) if the parent was the last entry.

Example (ls mode)

```
/
|-- [----a] HELLO.TXT  size=12          (2025-06-15 14:30:00)
|-- [----a] README.TXT size=11          (2025-06-15 14:30:00)
\-- [d----] SUBDIR/    (2025-06-15 14:30:00)
    |-- [----a] NESTED.TXT size=12        (2025-06-15 14:30:00)
    \-- [d----] DEEP/    (2025-06-15 14:30:00)
        \-- [----a] DEEP.TXT size=10      (2025-06-15 14:30:00)

Total: 4 file(s), 45 byte(s)
```

Example (cat mode)

```
$ ./fat16_tree test.img /SUBDIR/DEEP/DEEP.TXT
deep file
```

1.7 Exercise 5: File Writer

Task

Write `fat16_write.c` - the capstone exercise. Create a new file in the root directory of a FAT16 image. This requires modifying the FAT to allocate clusters, writing data to the data region, creating a new directory entry, and flushing everything back to the image file.

Usage

```
echo "hello world" | ./fat16_write test.img write HELLO.TXT 2025-06-15
14:30:00
./fat16_write test.img ls          # verify the listing
./fat16_write test.img cat HELLO.TXT # verify the content
```

What to implement (new)

- `find_free_cluster(fs, total_clusters)` - scan the FAT from cluster 2 for the first entry that equals `FAT16_FREE (0)`.
- `allocate_clusters(fs, file_size, total_clusters, &first)` - compute how many clusters are needed ($\lceil \text{file_size} \div \text{cluster_size} \rceil$), allocate them by updating the in-memory FAT, and link them into a chain terminated by `FAT16_EOC`.

- `write_data(fs, first_cluster, data, size)` - walk the allocated chain and write the content bytes to each cluster's location on disk.
- `create_dir_entry(root, count, filename, first_clus, size)` - find a free slot in the root directory (first byte 0x00 or 0xE5), fill in the 8.3 name, attributes (`ATTR_ARCHIVE`), first cluster, file size, and current timestamp.
- `flush_fat(fs)` - write the modified in-memory FAT back to *all* FAT copies in the image.
- `flush_root(fs, root, count)` - write the modified root directory back to the image.
- `fat16_write_file(fs, filename, data, size)` - orchestrate all the above steps.



Alert!

For this exercise, open the image in "`r+b`" mode (read *and* write), not "`rb`". Modify your `fat16_open()` accordingly.

Error handling

- File already exists: print Error: '`<name>`' already exists to stdout, exit 0.
- Disk full (no free clusters): print Error: `disk full` to stdout, exit 0.
- Root directory full (no free entry slot): print Error: `root directory full` to stdout, exit 0.
- Invalid filename (too long or empty): print Error: `invalid filename '<name>'` to stdout, exit 0.

Write procedure summary

1. Read content from `stdin` into a buffer.
2. Read the root directory from the image.
3. Verify the filename does not already exist.
4. Allocate clusters by scanning the FAT.
5. Write the data to the allocated clusters.
6. Create a directory entry in the root.
7. Flush the FAT and root directory back to the image.

Provided helpers

The template includes three helper functions you do **not** need to implement:

- `make_83_name(filename, name8, ext3)` - converts a filename like "HELLO.TXT" to the space-padded 8.3 format used in directory entries, returning 0 on success or -1 on invalid input.
- `encode_fat_date(year, month, day)` - packs a date into a 16-bit FAT date word.
- `encode_fat_time(hour, min, sec)` - packs a time into a 16-bit FAT time word.

1.8 Appendix: Creating Test Images

You can create a FAT16 test image on any Linux system using dosfstools:

```
# Create a 32 MB image filled with zeros
dd if=/dev/zero of=test.img bs=512 count=65536

# Format as FAT16 with a volume label
mkfs.fat -F 16 -n TESTDISK test.img

# Mount the image (requires root)
mkdir -p mnt
sudo mount -o loop test.img mnt

# Add some test files
echo "hello world" | sudo tee mnt/HELLO.TXT > /dev/null
echo "readme text" | sudo tee mnt/README.TXT > /dev/null
sudo mkdir -p mnt/SUBDIR/DEEP
echo "nested file" | sudo tee mnt/SUBDIR/NESTED.TXT > /dev/null
echo "deep file" | sudo tee mnt/SUBDIR/DEEP/DEEP.TXT > /dev/null

# Unmount
sudo umount mnt
```

Alternatively, on systems without root access, use mtools:

```
dd if=/dev/zero of=test.img bs=512 count=65536
mkfs.fat -F 16 -n TESTDISK test.img
# mtools operate directly on the image, no mount needed
echo "hello world" | mcopy -i test.img - ::HELLO.TXT
echo "readme text" | mcopy -i test.img - ::README.TXT
mmd -i test.img ::SUBDIR
echo "nested file" | mcopy -i test.img - ::SUBDIR/NESTED.TXT
```



2. Part II: Theoretical Exercises (Brightspace)

Submit your answers as a PDF via Brightspace. A L^AT_EX template (`os_students.tex`) is provided in the downloadable archive. Additionally, in order to differentiate between hand-written and large language model assisted code, we ask that you prefix your function names with the word `util` followed by an underscore, if they are model assisted. This is standard procedure in the course and is nothing to be concerned about.

Question 1 (4 points). A computer uses a programmable clock in square-wave mode. If a 500 MHz crystal is used, what should be the value of the holding register to achieve a clock resolution of

- (a) a millisecond (a clock tick once every millisecond)?
- (b) 100 microseconds?

Question 2 (4 points). In the discussion on journaling, it was shown that the disk can be recovered to a consistent state (a write either completes or does not take place at all) if a CPU crash occurs during a write. Does this property hold if the CPU crashes again during a recovery procedure. Explain your answer.

Question 3 (4 points). Compare RAID level 0 through 5 with respect to read performance, write performance, space overhead, and reliability.

Question 4 (4 points). CPU architects know that operating system writers hate imprecise interrupts. One way to please the OS folks is for the CPU to stop issuing new instructions when an interrupt is signaled, but allow all the instructions currently being executed to finish, then force the interrupt. Does this approach have any disadvantages? Explain your answer.

Question 5 (4 points). One mode that some DMA controllers use is to have the device controller send the word to the DMA controller, which then issues a second bus request to write to memory. How can this mode be used to perform memory to memory copy? Discuss any advantage or disadvantage of using this method instead of using the CPU to perform memory to memory copy.